

Christian B. Mendl, Richard M. Milbradt

Tutorial 4 (Canonical Polyadic (CP) decomposition for optimized matrix-matrix multiplication)
Using the outer product of vectors, the SVD of an $n \times n$ matrix $A = USV^\dagger$ can be represented as

$$A = \sum_{j=1}^n \sigma_j u_j \circ v_j^*,$$

where u_j and v_j are the column vectors of U and V , respectively. Note that the summation only needs to include the non-zero singular values, i.e., run from 1 to $r = \text{rank}(A)$. Similarly, a tensor $T \in \mathbb{C}^{n_1 \times n_2 \times n_3}$ can be decomposed as

$$T = \sum_{j=1}^r u_j \circ v_j \circ w_j =: \llbracket U, V, W \rrbracket,$$

with $U \in \mathbb{C}^{n_1 \times r}$, $V \in \mathbb{C}^{n_2 \times r}$, $W \in \mathbb{C}^{n_3 \times r}$ (not necessarily isometries). This is denoted the Canonical Polyadic (CP) decomposition. The minimal possible r defines the rank of the tensor.

(a) Does the CP decomposition always exist?

One usage of the CP decomposition is optimizing matrix-matrix multiplication, $C = AB$, with $A, B \in \mathbb{C}^{n \times n}$. A literal implementation of $c_{ik} = \sum_{j=1}^n a_{ij} b_{jk}$ (for $i, k = 1, \dots, n$) has runtime $\mathcal{O}(n^3)$, but it turns out that this can be further improved. In the following, we assume that $n = 2^k$ is a power of 2, and use a recursive block partitioning into four blocks of size $2^{k-1} \times 2^{k-1}$ each, such that

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}.$$

- (b) Working with this block partitioning, write out the algorithmic complexity of the naive matrix multiplication. What is the computationally most expensive step in this algorithm?
- (c) We can enumerate the block indexing of the submatrices as $(1, 1) \rightarrow 1$, $(1, 2) \rightarrow 2$, $(2, 1) \rightarrow 3$, $(2, 2) \rightarrow 4$, allowing for a one-hot encoding e_1 to e_4 . Based on that, assemble a tensor of degree 3, with the first two dimensions corresponding to the input block indices and the third to the output index.
- (d) How does the rank of this tensor affect the complexity of the matrix-matrix multiplication?

Remark: It turns out that the tensor in (c) has rank 7. This leads to the *Strassen algorithm*, exploiting that only 7 block matrix multiplications are necessary at each level of recursion to perform the same operation. The algorithmic complexity is thus reduced from $\mathcal{O}(n^3)$ to $\mathcal{O}(7^k) = \mathcal{O}(n^{\log_2 7}) \approx \mathcal{O}(n^{2.80735})$.¹

Solution

- (a) Yes, this is possible. You can easily replace all three vectors with basis vectors multiplied by a constant.
- (b) The naive algorithm is

$$\begin{aligned} C_{11} &= A_{11}B_{11} + A_{12}B_{21} \\ C_{12} &= A_{11}B_{12} + A_{12}B_{22} \\ C_{21} &= A_{21}B_{11} + A_{22}B_{21} \\ C_{22} &= A_{21}B_{12} + A_{22}B_{22} \end{aligned}$$

This requires $8 \left(\frac{n}{2}\right)^3 = \mathcal{O}(n^3)$ operations. The most expensive step is the matrix multiplication, which takes $\mathcal{O}(m^3)$ when multiplying two $m \times m$ matrixes.

¹The up to date smallest-known exponent is $\mathcal{O}(n^{2.37286})$, see arxiv.org/abs/2010.05846.

(c) We can encode the 8 different matrix multiplications as follows:

$$C_{ik} += A_{ij}B_{jk} \rightarrow e_{ij} \circ e_{jk} \circ e_{ik}$$

This gives for the expression in part (b):

$$\begin{aligned}\mathcal{T} = & e_{11} \circ e_{11} \circ e_{11} + e_{12} \circ e_{21} \circ e_{11} \\ & + e_{11} \circ e_{12} \circ e_{12} + e_{12} \circ e_{22} \circ e_{12} \\ & + e_{21} \circ e_{11} \circ e_{21} + e_{22} \circ e_{21} \circ e_{21} \\ & + e_{21} \circ e_{12} \circ e_{22} + e_{22} \circ e_{22} \circ e_{22}.\end{aligned}$$

(d) This tensor has rank less than 8 (actually rank 7), according to the following CP decomposition with 7 terms:

$$\begin{aligned}\mathcal{T} = & (e_{11} + e_{22}) \circ (e_{11} + e_{22}) \circ (e_{11} + e_{22}) \\ & + (e_{21} + e_{22}) \circ e_{11} \circ (e_{21} - e_{22}) \\ & + e_{11} \circ (e_{12} - e_{22}) \circ (e_{12} + e_{22}) \\ & + e_{22} \circ (e_{21} - e_{11}) \circ (e_{11} + e_{21}) \\ & + (e_{11} + e_{12}) \circ e_{22} \circ (-e_{11} + e_{12}) \\ & + (e_{21} - e_{11}) \circ (e_{11} + e_{12}) \circ e_{22} \\ & + (e_{12} - e_{22}) \circ (e_{21} + e_{22}) \circ e_{11}.\end{aligned}$$

Hence, we can then compute the matrix multiplication as follows:

$$\begin{aligned}M_1 &= (A_{11} + A_{22})(B_{11} + B_{22}) \\ M_2 &= (A_{21} + A_{22})B_{11} \\ M_3 &= A_{11}(B_{12} - B_{22}) \\ M_4 &= A_{22}(-B_{11} + B_{21}) \\ M_5 &= (A_{11} + A_{22})B_{22} \\ M_6 &= (-A_{11} + A_{21})(B_{11} + B_{12}) \\ M_7 &= (A_{12} - A_{22})(B_{21} + B_{22})\end{aligned}$$

The entries of C are then

$$\begin{aligned}C_{11} &= M_1 + M_4 - M_5 + M_7 \\ C_{12} &= M_3 + M_5 \\ C_{11} &= M_2 + M_4 \\ C_{11} &= M_1 - M_2 + M_3 + M_6\end{aligned}$$

The algorithmic complexity is thus reduced from $\mathcal{O}(n^3)$ to $\mathcal{O}(7^k) = \mathcal{O}(n^{\log_2 7}) \approx \mathcal{O}(n^{2.80735})$ when the shown steps are performed recursively.